

36-315: Statistical Graphics & Visualization

- course objectives
 - learn useful principles for making appropriate statistical graphics.
 - critique existing graphs and remake better ones.
 - visualize statistical analyses to facilitate communication.
 - pinpoint the statistical claims you can/cannot make from graphics.
 - write and speak publicly about statistical graphics.
 - practice tidy data manipulation in R using the tidyverse
 - practice reproducible workflows with Quarto
- grammar of graphics defined and used in ggplot2
 - see 01lec.pdf
- goal of data visualization: show data, communicate a story
 - induce viewer to think about substance, not graphical methodology
 - make large, complex datasets more coherent
 - encourage comparison of different pieces of data
 - describe, explore, and identify relationships
 - avoid data distortion and data decoration
 - use consistent graph design
 - avoid graphs that lead to misleading conclusions!
- data types
 - quantitative
 - discrete
 - continuous
 - categorical (factor)
 - nominal
 - no order
 - e.g. race, species
 - ordinal
 - ordered!
 - ranking
 - DEFAULT IN R! manually define factor levels, or alpha. default
- area plots
 - pie chart (BAD!!!)
 - bar chart
 - stacked bar/spine chart (for variable comparison)
 - waffle charts????
 - rose diagrams (temporal or directional context can justify usage)
- something something “geom_bar(stat=identity)” to “take y as is”
- α level CI is $\hat{x} \pm z_{1-\frac{\alpha}{2}} \cdot SE(\hat{x})$
- 1d chisq test: H_A is “at least one category differs”
`chisq.test(table(penguins$species))`
- CI interpretation
 - If CIs don't overlap → significant difference
 - If CIs overlap → a little ambiguous
 - If CIs overlap → a lot no significant difference

- multiple testing:
 - have multiple pairwise comparisons via CI eyeballing? Type 1 error is now above 5%!!!
 - correct by inflating p values
 - Bonferroni Correction:
 - making K comparisons \rightarrow reject iff $p \leq \frac{\alpha}{K}$.
 - easy to impl and popular but inflates p the most
 - CIs: plot $(1 - \alpha)\%$ CIs \rightarrow plot $(1 - \frac{\alpha}{K})\%$ CIs
- 2d chisq test: H_A is “ A, B independent”


```
chisq.test(table(penguins$species, penguins$island))
```

 - visualize this with mosaic plots
- mosaic plots: can shade by Pearson residuals
 - more positive p.r \rightarrow more counts than expected, more neg is vice versa
 - we might reject null for the global chisq test but see all white residuals: can't reject null for individual local tests.
- 1d quant
 - boxplots: only summary stats: bad!
 - hist: see dist, bin width matters.
 - density curves: conditional dists
- estimation schools of thought:
 - parametric: assume dist, est params (eg MLE, 3623X)
 - nonparametric: make few assumptions, use whole dataset (density curves, regression lines??)
- kernel density estimation
 - place lil dist on every x_i
 - usually normal, but many exist (fuck it triangle. things can help maintain strict left right dist bounds if needed)
 - bandwidth (higher is more smooth dist): ggplot alr uses Gaussian reference rule of thumb, set to $1.06 \cdot \sigma_{\text{sample}} \cdot n^{-1/5}$.
 - adjust bandwith via `geom_density(adjust = <multiplier>)`